

Examples for large sample credible intervals

Example 1: Prevalence of Recessive Gene

If gene frequencies are in equilibrium, the genotypes AA , Aa , and aa occur with probabilities $(1-\theta)^2$, $2\theta(1-\theta)$, and θ^2 respectively, where θ represents the overall prevalence of the recessive a gene in the population. Plato et al. (1964) published the following data on a haptoglobin type in a sample of 190 people:

Haptoglobin Type	AA	Aa	aa
Count	112	68	10

Let's regard the vector $\mathbf{x} = (x_1, x_2, x_3) = (112, 68, 10)$ as a realization of the random variable $\mathbf{X} \sim \text{Multinomial}((1-\theta)^2, 2\theta(1-\theta), \theta^2)$.

To save some time/allow us to focus on the results of interest here, I'll give you the likelihood function, its first and second derivatives with respect to θ , and the form of the posterior:

Preliminary Results

General form of Multinomial pmf

$$f(\mathbf{x}|\mathbf{p}) = \frac{n!}{x_1!x_2!\cdots x_k!} p_1^{x_1} p_2^{x_2} \cdots p_k^{x_k}$$

Likelihood function

In our example, $p_1 = (1-\theta)^2$, $p_2 = 2\theta(1-\theta)$, and $p_3 = \theta^2$.

$$\begin{aligned}\mathcal{L}(\theta|\mathbf{x}) &= f(\mathbf{x}|\theta) \\ &= \{(1-\theta)^2\}^{x_1} \{2\theta(1-\theta)\}^{x_2} \{\theta^2\}^{x_3}\end{aligned}$$

I'm going to leave this in terms of x_1, x_2 , and x_3 for now.

Log-likelihood function

$$\begin{aligned}\ell(\theta|\mathbf{x}) &= \log[\mathcal{L}(\theta|\mathbf{x})] \\ &= \log [\{(1-\theta)^2\}^{x_1} \{2\theta(1-\theta)\}^{x_2} \{\theta^2\}^{x_3}] \\ &= x_1 \log\{(1-\theta)^2\} + x_2 \log\{2\theta(1-\theta)\} + x_3 \log\{\theta^2\}\end{aligned}$$

First and second derivatives of log-likelihood function

The first derivative of the log-likelihood is:

$$\frac{d}{d\theta} \ell(\theta|\mathbf{x}) = \cdots = \frac{-2x_1\theta}{\theta(1-\theta)} + \frac{x_2(1-2\theta)}{\theta(1-\theta)} + \frac{2x_3(1-\theta)}{\theta(1-\theta)}$$

The second derivative of the log-likelihood is:

$$\frac{d^2}{d\theta^2}\ell(\theta|\mathbf{x}) = \dots = -\frac{2x_1 + x_2}{(1-\theta)^2} - \frac{2x_3 + x_2}{\theta^2}$$

Maximum likelihood estimator

Setting the first derivative equal to 0, we obtain a maximum likelihood estimator of

$$\hat{\theta}^{MLE} = \frac{X_2 + 2X_3}{2n}.$$

It can be verified that this gives a global maximum of the likelihood function.

Posterior Distribution

Suppose we adopt a prior of $\Theta \sim \text{Uniform}(0, 1)$

The prior distribution for Θ has density $f_{\Theta}(\theta) = \begin{cases} 1 & \text{if } \theta \in [0, 1] \\ 0 & \text{otherwise} \end{cases}$.

Additionally, in part (a) we showed that $f_{\mathbf{X}|\Theta}(\mathbf{x}|\theta) = \{(1-\theta)^2\}^{x_1} \{2\theta(1-\theta)\}^{x_2} \{\theta^2\}^{x_3}$.

Applying Bayes' Rule, we find that

$$f_{\Theta|\mathbf{X}}(\theta|\mathbf{x}) = \dots = \begin{cases} c\{(1-\theta)^2\}^{x_1} \{2\theta(1-\theta)\}^{x_2} \{\theta^2\}^{x_3} & \text{if } \theta \in [0, 1] \\ 0 & \text{otherwise} \end{cases}$$

The integral is kindof annoying, but can be done.

Problems for you

1. Find a large-sample normal approximation to the posterior distribution for θ .

You should find the numeric values of the mean and variance of this normal approximation based on our data set.

Mean:

```
mle <- (68 + 2 * 10)/(2 * 190)
mle
```

```
## [1] 0.2315789
```

Variance:

```
second_deriv_loglik <- -(2 * 112 + 68)/(1 - mle)^2 - (2 * 10 + 68)/(mle^2)
post_approx_var <- -1/second_deriv_loglik
post_approx_var
```

```
## [1] 0.0004682898
```

Therefore, the posterior distribution for θ is approximately Normal(0.232, 0.000468)

2. Add a plot of the pdf of the normal approximation to the plot of the actual posterior below.


I used Wolfram Alpha to figure out how to calculate the constant c . That's the first 4 lines of the calculation of the log density in the `dposterior` function below:

```
library(ggplot2)

dposterior <- function(theta, x_1, x_2, x_3, log = FALSE) {
  n <- x_1 + x_2 + x_3
  log_d_posterior <- sum(log(seq_len(2 * n + 1))) -
    x_2 * log(2) -
    sum(log(seq_len(2 * x_1 + x_2))) -
    sum(log(seq_len(x_2 + 2 * x_3))) +
    2 * x_1 * log(1 - theta) +
    x_2 * log(2 * theta * (1 - theta)) +
    2 * x_3 * log(theta)

  if(log) {
    return(log_d_posterior)
  } else {
    return(exp(log_d_posterior))
  }
}

ggplot(data = data.frame(theta = c(0, 1)), mapping = aes(x = theta)) +
  stat_function(fun = dposterior, args = list(x_1 = 112, x_2 = 68, x_3 = 10), n = 1001) +
  stat_function(fun = dnorm, args = list(mean = 0.232, sd = sqrt(0.000468)), color = "orange")
```



3. Find and interpret an approximate Bayesian 95% credible interval for θ based on the normal approximation.

```
qnorm(c(0.025, 0.975), mean = 0.232, sd = sqrt(0.000468))
```

```
## [1] 0.1895995 0.2744005
```

There is probability 0.95 that the proportion of people in the population carrying the recessive gene is between about 0.19 and 0.27.